

Double-click (or enter) to edit

```
1 import pandas as pd
```


```
1 import pandas as pd
2 data = pd.read_csv('/content/creditcard.csv')
3 data = data.dropna() # Example for handling missing values
4 data.head()
```




	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	
0	0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0
1	0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0
2	1	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0
3	1	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1
4	2	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0

5 rows × 31 columns


Distributions




2-d distributions



Time series



Values



```
1 from matplotlib import pyplot as plt
2 _df_7.plot(kind='scatter', x='V3', y='V4', s=32, alpha=.8)
3 plt.gca().spines[['top', 'right',]].set_visible(False)
```



```
-----
NameError                                Traceback (most recent call last)
<ipython-input-1-8d17c3903423> in <cell line: 2>()
      1 from matplotlib import pyplot as plt
----> 2 _df_7.plot(kind='scatter', x='V3', y='V4', s=32, alpha=.8)
      3 plt.gca().spines[['top', 'right',]].set_visible(False)

NameError: name '_df_7' is not defined
```

```
1 from matplotlib import pyplot as plt
2 import seaborn as sns
3 def _plot_series(series, series_name, series_index=0):
4     palette = list(sns.palettes.mpl_palette('Dark2'))
5     xs = series['Time']
6     ys = series['V3']
7
8     plt.plot(xs, ys, label=series_name, color=palette[series_index % len(palette)])
9
10 fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
11 df_sorted = _df_10.sort_values('Time', ascending=True)
```

```

12 _plot_series(df_sorted, '')
13 sns.despine(fig=fig, ax=ax)
14 plt.xlabel('Time')
15 _ = plt.ylabel('V3')

```

```

1 from matplotlib import pyplot as plt
2 _df_6.plot(kind='scatter', x='V2', y='V3', s=32, alpha=.8)
3 plt.gca().spines[['top', 'right']].set_visible(False)

```

```

1 from matplotlib import pyplot as plt
2 _df_3['V3'].plot(kind='hist', bins=20, title='V3')
3 plt.gca().spines[['top', 'right']].set_visible(False)

```

Balance the Dataset: Fraudulent transactions are typically a small percentage of the data. Use techniques like oversampling the minority class (fraud) or undersampling the majority class (non-fraud).

```

1 # prompt: Balance the Dataset: Fraudulent transactions are typically a small percentage of the data. Use techniques like oversampling the
2
3 import pandas as pd
4 from sklearn.model_selection import train_test_split
5
6 # Separate the data into fraudulent and non-fraudulent transactions
7 fraudulent_transactions = data[data['Class'] == 1]
8 non_fraudulent_transactions = data[data['Class'] == 0]
9
10 # Oversample the fraudulent transactions
11 oversampled_fraudulent_transactions = fraudulent_transactions.sample(
12     len(non_fraudulent_transactions), replace=True)
13
14 # Combine the oversampled fraudulent transactions with the non-fraudulent transactions
15 balanced_data = pd.concat([oversampled_fraudulent_transactions, non_fraudulent_transactions])
16
17 # Split the balanced data into training and testing sets
18 X_train, X_test, y_train, y_test = train_test_split(
19     balanced_data.drop('Class', axis=1), balanced_data['Class'], test_size=0.2, random_state=42)
20

```

Extract and engineer relevant features that can help in fraud detection, such as transaction frequency, average transaction amount, etc.

```

1 # prompt: Extract and engineer relevant features that can help in fraud detection, such as transaction frequency, average transaction amc
2
3 import pandas as pd
4 # Transaction frequency per user
5 user_transaction_counts = data.groupby('V1')['Class'].count()
6
7 # Average transaction amount per user
8 user_average_amounts = data.groupby('V1')['Amount'].mean()
9
10 # Combine the extracted features into a new DataFrame
11 engineered_features = pd.DataFrame({
12     'Transaction_Frequency': user_transaction_counts,
13     'Average_Amount': user_average_amounts
14 })
15
16 # Merge the engineered features with the original data
17 data = data.merge(engineered_features, on='V1', how='left')
18

```

Exploratory Data Analysis (EDA):

```

1 # prompt: Exploratory Data Analysis (EDA):
2
3 import pandas as pd
4
5 # Load the data
6 data = pd.read_csv('/content/creditcard.csv')
7
8 # Check the shape of the data
9 print(data.shape)
10
11 # Check the head of the data
12 print(data.head())
13
14 # Check the tail of the data
15 print(data.tail())
16

```

```

17 # Check the data types of the data
18 print(data.dtypes)
19
20 # Check the descriptive statistics of the data
21 print(data.describe())
22
23 # Check the correlation matrix of the data
24 print(data.corr())
25
26 # Check the distribution of the data
27 import matplotlib.pyplot as plt
28
29 data.hist(figsize=(10, 10))
30 plt.show()
31
32 # Check the boxplots of the data
33 data.boxplot(figsize=(10, 10))
34 plt.show()
35

```

```

(114962, 31)
Time      V1      V2      V3      V4      V5      V6      V7  \
0      0 -1.359807 -0.072781  2.536347  1.378155 -0.338321  0.462388  0.239599
1      0  1.191857  0.266151  0.166480  0.448154  0.060018 -0.082361 -0.078803
2      1 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499  0.791461
3      1 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609
4      2 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941

      V8      V9      ...      V21      V22      V23      V24      V25  \
0  0.098698  0.363787  ... -0.018307  0.277838 -0.110474  0.066928  0.128539
1  0.085102 -0.255425  ... -0.225775 -0.638672  0.101288 -0.339846  0.167170
2  0.247676 -1.514654  ...  0.247998  0.771679  0.909412 -0.689281 -0.327642
3  0.377436 -1.387024  ... -0.108300  0.005274 -0.190321 -1.175575  0.647376
4 -0.270533  0.817739  ... -0.009431  0.798278 -0.137458  0.141267 -0.206010

      V26      V27      V28  Amount  Class
0 -0.189115  0.133558 -0.021053   149.62   0.0
1  0.125895 -0.008983  0.014724    2.69   0.0
2 -0.139097 -0.055353 -0.059752   378.66   0.0
3 -0.221929  0.062723  0.061458   123.50   0.0
4  0.502292  0.219422  0.215153    69.99   0.0

[5 rows x 31 columns]
Time      V1      V2      V3      V4      V5      V6  \
114957  73690 -2.197480  1.982629  0.259502  0.924323 -0.879938 -0.135952
114958  73690 -2.197480  1.982629  0.259502  0.924323 -0.879938 -0.135952
114959  73690  1.255655  0.293362  0.288616  0.701727 -0.447134 -1.093442
114960  73690  1.270638 -0.089535 -0.990690 -0.375303  0.652307 -0.073908
114961  73691  1.295784  0.049457 -1.301814 -0.476648  2.131357  3.184446

      V7      V8      V9      ...      V21      V22      V23  \
114957 -0.380186  0.611134  0.278764  ... -0.033516 -0.367199 -0.099698
114958 -0.380186  0.611134  0.278764  ... -0.033516 -0.367199 -0.099698
114959  0.029565 -0.155947  0.160172  ... -0.297287 -0.898954  0.139494
114960  0.394543 -0.164288 -0.317251  ... -0.067492 -0.344301 -0.372630
114961 -0.494416  0.802781 -0.075014  ...      NaN      NaN      NaN

      V24      V25      V26      V27      V28  Amount  Class
114957 -0.111166 -0.182825  0.427261 -0.895134  0.164611    9.51   0.0
114958 -0.111166 -0.182825  0.427261 -0.895134  0.164611    9.51   0.0
114959  0.322281  0.187536  0.097228 -0.028586  0.029014    1.79   0.0
114960 -1.296908  0.731021  1.152123 -0.131651 -0.024591   75.00   0.0
114961      NaN      NaN      NaN      NaN      NaN      NaN   NaN

[5 rows x 31 columns]
Time      int64
V1      float64
V2      float64
V3      float64
V4      float64
V5      float64
V6      float64
V7      float64
V8      float64
V9      float64
V10     float64
V11     float64
V12     float64

```

Choose classification algorithms such as Random Forest, Support Vector Machines (SVM), or Neural Networks.

```

V15      float64

1 models = {
2     "Random Forest": RandomForestClassifier(),
3     "Support Vector Machine": SVC(),
4     "Neural Network": MLPClassifier(max_iter=1000)
5 }

```

```
6
7 # Splitting the dataset
8
9 # Features and target
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
11
12 # Importing the Random Forest Classifier
13 from sklearn.ensemble import RandomForestClassifier
14
15 # Creating the model
16 model = RandomForestClassifier()
17
18 # Training the model
19 model.fit(X_train, y_train)
20
21 # Making predictions
22 y_pred = model.predict(X_test)
23
24 # Evaluating the model
25 print('Accuracy:', accuracy_score(y_test, y_pred))
26 print('Precision:', precision_score(y_test, y_pred))
27 print('Recall:', recall_score(y_test, y_pred))
28 print('F1-Score:', f1_score(y_test, y_pred))
29
30 # Output
31 Accuracy: 0.9999346234309623
32 Precision: 0.9998689040377556
33 Recall: 1.0
34 F1-Score: 0.9999344477220584
```


HERE TRAINING THE MODEL

```
1 # Importing the necessary libraries
2 import numpy as np
3 import pandas as pd
4 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
5 from sklearn.model_selection import train_test_split
6 from sklearn.ensemble import RandomForestClassifier
7
8 # Load the data
9 data = pd.read_csv('/content/creditcard.csv')
10
11 # Separate the data into fraudulent and non-fraudulent transactions
12 fraudulent_transactions = data[data['Class'] == 1]
13 non_fraudulent_transactions = data[data['Class'] == 0]
14
15 # Oversample the fraudulent transactions
16 oversampled_fraudulent_transactions = fraudulent_transactions.sample(
17     len(non_fraudulent_transactions), replace=True)
18
19 # Combine the oversampled fraudulent transactions with the non-fraudulent transactions
20 balanced_data = pd.concat([oversampled_fraudulent_transactions, non_fraudulent_transactions])
21
22 # Split the balanced data into training and testing sets
23 X_train, X_test, y_train, y_test = train_test_split(
24     balanced_data.drop('Class', axis=1), balanced_data['Class'], test_size=0.2, random_state=42)
25
26 # Train the Random Forest model
27 model = RandomForestClassifier()
28 model.fit(X_train, y_train)
29
30 # Make predictions on the test set
31 y_pred = model.predict(X_test)
32
33 # Evaluate the model
34 print("Accuracy:", accuracy_score(y_test, y_pred))
35 print("Precision:", precision_score(y_test, y_pred))
36 print("Recall:", recall_score(y_test, y_pred))
37 print("F1-Score:", f1_score(y_test, y_pred))
38
39 # Output
40 Accuracy: 0.9999346234309623
41 Precision: 0.9998689040377556
42 Recall: 1.0
43 F1-Score: 0.9999344477220584
```



FNIAL TEST


```
1 # Visualization of results
2 metrics = {
3     'Accuracy': accuracy_score(y_test, y_pred),
4     'Precision': precision_score(y_test, y_pred),
5     'Recall': recall_score(y_test, y_pred),
6     'F1-Score': f1_score(y_test, y_pred)
7 }
8
9 metrics_df = pd.DataFrame.from_dict(metrics, orient='index', columns=['Score'])
10 metrics_df.plot(kind='bar', legend=False)
11 plt.title('Model Evaluation Metrics')
12 plt.ylabel('Score')
13 plt.show()
14
```



```
1 # prompt: GIVE THE INFORMATION ABOVE THE RESULT
2
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 # Print the results
6 print("Accuracy:", accuracy_score(y_test, y_pred))
7 print("Precision:", precision_score(y_test, y_pred))
8 print("Recall:", recall_score(y_test, y_pred))
9 print("F1-Score:", f1_score(y_test, y_pred))
10
11 # Visualization of results
12 metrics = {
13     'Accuracy': accuracy_score(y_test, y_pred),
14     'Precision': precision_score(y_test, y_pred),
15     'Recall': recall_score(y_test, y_pred),
16     'F1-Score': f1_score(y_test, y_pred)
17 }
18
19 metrics_df = pd.DataFrame.from_dict(metrics, orient='index', columns=['Score'])
20 metrics_df.plot(kind='bar', legend=False)
21 plt.title('Model Evaluation Metrics')
22 plt.ylabel('Score')
23 plt.show()
24
```



```
Accuracy: 0.9999346234309623
Precision: 0.9998689040377556
Recall: 1.0
F1-Score: 0.9999344477220584
```



20000

Explanation of Metrics: Accuracy:

Definition: The ratio of correctly predicted instances to the total instances. Interpretation: An accuracy of 0.9999 means that 99.99% of the transactions were correctly classified as either fraudulent or non-fraudulent. Precision:

Definition: The ratio of correctly predicted positive observations to the total predicted positives. Interpretation: A precision of 0.9999 indicates that 99.99% of the transactions identified as fraud were indeed fraud. Recall (Sensitivity or True Positive Rate):

Definition: The ratio of correctly predicted positive observations to all observations in the actual class. Interpretation: A recall of 1.0 means that the model correctly identified all actual fraud cases in the dataset. F1-Score:

Definition: The weighted average of Precision and Recall. The F1 Score takes both false positives and false negatives into account. Interpretation: An F1-Score of 0.9999 indicates a balanced performance of the model in terms of both precision and recall. Visualization: The bar chart displays these metrics for a visual comparison. Each bar represents the score for a different metric, allowing for an easy comparison of model performance across these important evaluation criteria.

Potential Overfitting: The scores are very high (almost perfect), which might suggest that the model could be overfitting the training data, especially if the dataset is not large or diverse enough. Overfitting occurs when a model learns the training data too well, including noise and outliers, and performs poorly on unseen data. Next Steps: Cross-Validation: To ensure the model generalizes well, consider using cross-validation. Testing on Unseen Data: Evaluate the model on a separate test set that was not used during the training process. Model Comparison: Compare the performance of different models (Random Forest, SVM, Neural Network) to choose the best one. Feature Importance: Analyze feature importance to understand which features contribute most to fraud detection.

